

# Scalable Linear Solvers for GPU-accelerated heterogeneous supercomputers

Pasqua D'Ambra

Institute for Applied Computing, National Research Council (IAC-CNR)  
and CINI Lab on HPC-KTT

[pasqua.dambra@cnr.it](mailto:pasqua.dambra@cnr.it)

2024 Frontiers of High-Performance Computing  
in Modeling and Simulation  
July 4-5, 2024



# Greetings

## HPCTeam@IAC

- Massimo Bernaschi (IAC-CNR), IT
- Mauro G. Carrozzo (IAC-CNR), IT
- Alessandro Celestini (IAC-CNR), IT
- Fabio Durastante (Univ. of Pisa and IAC-CNR), IT
- Salvatore Filippone (Univ. of Rome Tor-Vergata and IAC-CNR), IT
- Giacomo Piperno (IAC-CNR), IT

## Collaborations

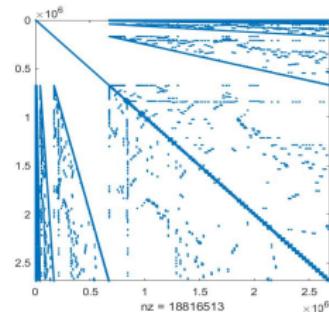
- Mahantesh M. Halappanavar and S M Ferdous , PNNL (Richland, WA), USA
- Stefano Massei, (Univ. of Pisa), IT
- Alex Pothen, Purdue University (West Lafayett, IN), USA
- Stephen Thomas, NREL, (Golden, CO), USA
- Panayot S. Vassilevski, Portland State University (Portland, OR), USA
- Ludmil Zikatanov, The Penn State University, PSU (State College), USA

# What we want to solve

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \text{ (s.p.d.)} \quad x, b \in \mathbb{R}^n$$

*n large*

$$\text{sparsity degree} = 1 - \frac{\text{nnz}}{n^2} \approx 1$$



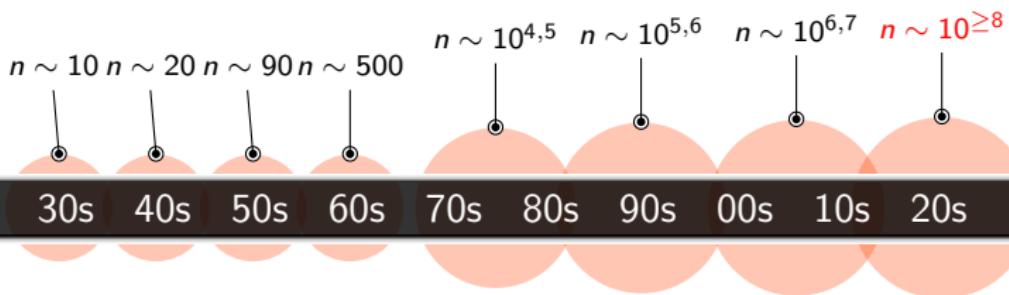
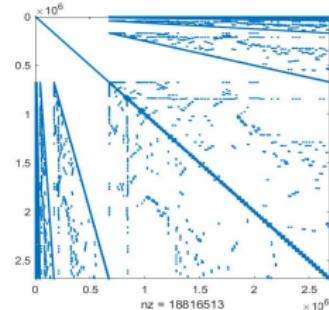
often the most time consuming computational kernel in many areas of Computational/Data Science

# What we want to solve

$$Ax = b, \quad A \in \mathbb{R}^{n \times n} \text{ (s.p.d.)} \quad x, b \in \mathbb{R}^n$$

**$n$  large**

$$\text{sparsity degree} = 1 - \frac{nnz}{n^2} \approx 1$$



The **exascale** challenge: using current supercomputers, targeting systems doing  $10^{18}$  Flops, to solve problems with **tens of trillions** ( $10^{13}$ ) dofs

# Energy oriented Center of Excellence (EuroHPC-JU EoCoE)

## Wind Models

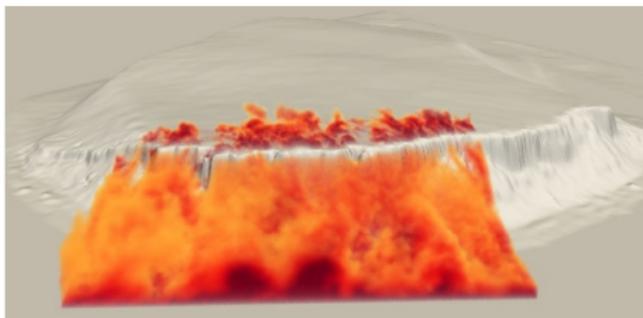
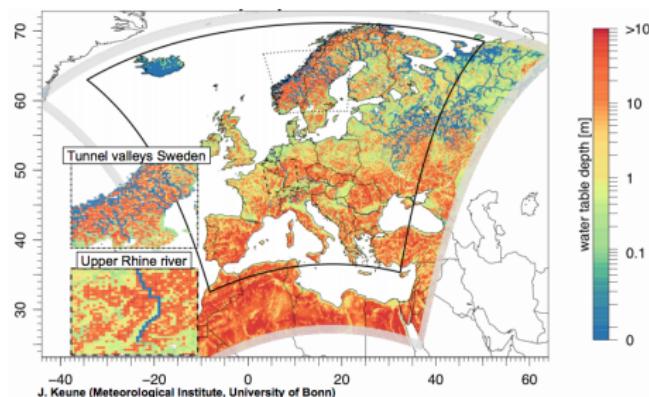


Image credits H. Owen and G. Marin, BSC

- Navier-Stokes equations,
- Euler equations,
- Turbulence models,
- ...

## Regional Hydrological Models



- Darcy equation,
- Richards equation,
- Equations for overland flow

Target dofs:  $n > 10^{13}$

# Where we want to run<sup>1</sup>

	System	Cores	Rmax (PFlops)	HPCG (PFlops)	
7	Leonardo	1,824,768	241.20	3.1	
8	MareNostrum 5 ACC	663,040	175.30	1.2	
21	Juwels Booster	449,280	44.12	1.3	Leonardo - Cineca

- Computers with thousands of CPU cores and GPU accelerators
- Deep memory hierarchies and hybrid forms of parallelism/programming models&tools (MPI, OpenMP, OpenACC, Cuda, ...)



Juwels - JSC

HPCG obtain a very small percentage ( $\approx 1.3\%$  on Leonardo) of the peak performance!

<sup>1</sup>TOP500 list, June 2024 – <https://www.top500.org>



# Parallel Sparse Computation Toolkit



recognized as “Excellent Science Innovation”  
by the EU Innovation Radar



# (Algebraic) MultiGrid methods

**V-cycle**( $I, nlev, A^I, b^I, x^I$ )

if ( $I \neq nlev$ ) then

$$x^I = x^I + (M^I)^{-1} (b^I - A^I x^I)$$

$$b^{I+1} = (P^I)^T (b^I - A^I x^I)$$

$$x^{I+1} = \text{V-cycle}(I + 1, A^{I+1}, b^{I+1}, 0)$$

$$x^I = x^I + P^I x^{I+1}$$

$$x^I = x^I + (M^I)^{-T} (b^I - A^I x^I)$$

else

$$x^I = (A^I)^{-1} b^I$$

endif

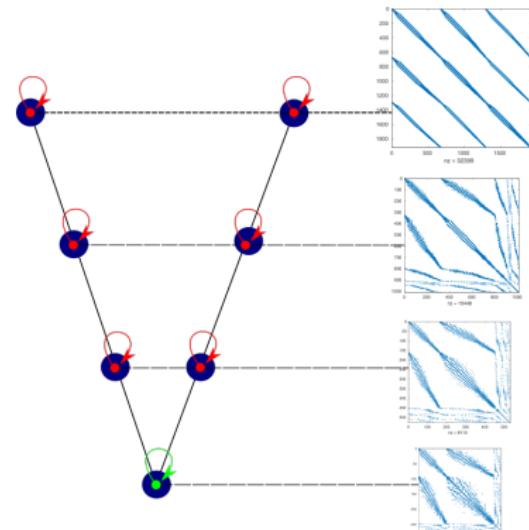
return  $x^I$

end

**Smoothes**

$$M^I : \mathbb{R}^{n_I} \rightarrow \mathbb{R}^{n_I}$$

“damping high frequencies”



**Prolongator**

$$P^I : \mathbb{R}^{n_{I+1}} \rightarrow \mathbb{R}^{n_I}$$

“transferring low frequencies”

# MultiGrid Convergence

Theorem (McCormick 1985, Vassilevski 2008)

If  $M^l$  is a contraction at each level  $l$ , i.e.,  $\|I - (M^l)^{-1}A^l\|_{A^l} < 1$ , the V-cycle preconditioner  $B$  defined as the multiplicative composition of the iteration matrix

$$I - (B^l)^{-1}A^l = (I - (M^l)^{-T}A^l)(I - P^l((P^l)^TA^lP^l)^{-1}(P^l)^TA^l)(I - (M^l)^{-1}A^l)$$

has the following error bound:

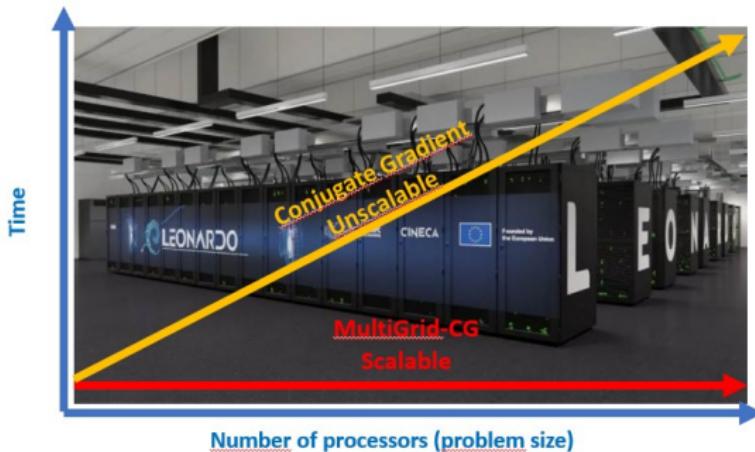
$$\|E\|_A^2 = \|I - B^{-1}A\|_A^2 \leq 1 - \frac{1}{C} \quad \text{with}$$
$$C = \max_l C^l$$

where  $C^l = \sup_{v \in \text{Range}(P^l)^\perp \setminus 0} \frac{\|v\|_{M^l}^2}{\|v\|_A^2} \geq 1$  is the strong approximation constant and  $\tilde{M}^l = M^l(M^l + (M^l)^T - A^l)^{-1}(M^l)^T$  is the symmetrized smoother.

Optimal Convergence (independent of problem size and number of levels)

# Scalable (AMG) preconditioners

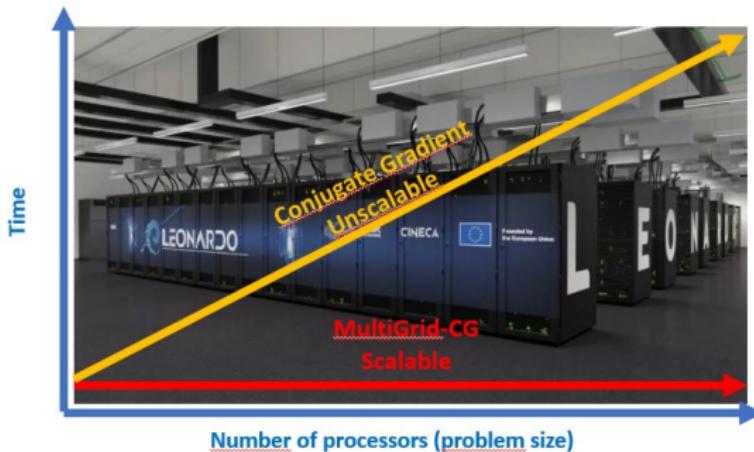
(inspired by Rob Falgout)



- Optimal complexity: AMG can be optimal ( $\mathcal{O}(nnz)$  flops), then have good scalability potential

# Scalable (AMG) preconditioners

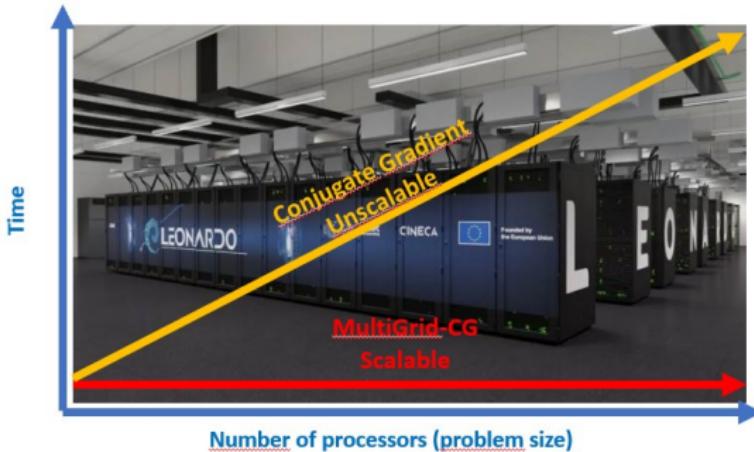
(inspired by Rob Falgout)



- Optimal complexity: AMG can be optimal ( $\mathcal{O}(nnz)$  flops), then have good scalability potential
- Optimal convergence:  $\|E\|_A^2 < 1$  being independent of  $nnz$  (Algorithmic scalability)

# Scalable (AMG) preconditioners

(inspired by Rob Falgout)



Optimal complexity and optimal convergence are not sufficient for HPC!

**Implementation scalability:**  $B$  should be composed of local actions essentially based on a “hierarchy” of sparse matrix-vector products

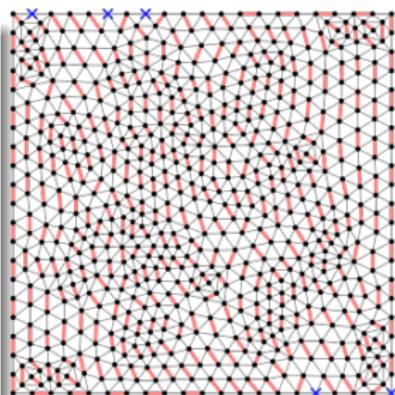
# Parallel AMG setup

## CMATCH algorithm

Given a graph  $G = (\mathcal{V}, \mathcal{E})$  (with adjacency matrix  $A$ ), and a weight (smooth) vector  $\mathbf{w}$  we consider the weighted version of  $G$  obtained by considering the weight matrix  $\hat{A}$ :

$$(\hat{A})_{i,j} = \hat{a}_{i,j} = 1 - \frac{2a_{i,j}w_i w_j}{a_{i,i}w_i^2 + a_{j,j}w_j^2},$$

- a *matching*  $\mathcal{M}$  is a set of pairwise non-adjacent edges, containing no loops;
  - a **maximum product matching** maximizes the product of the weights of its edges  $e_{i \rightarrow j}$ .
- 
- Increase coarsening ratio by performing **more than one sweep of matching** per level
  - Increase regularity of  $P_{l+1}^l$  with **smoothed prolongator** by applying one step of Jacobi method



Divide the index set into  
matched vertices

$$\mathcal{I} = \bigcup_{i=1}^{n_p} \mathcal{G}_i, \text{ with } \mathcal{G}_i \cap \mathcal{G}_j = \emptyset \text{ if } i \neq j,$$

# Parallel Smoothers

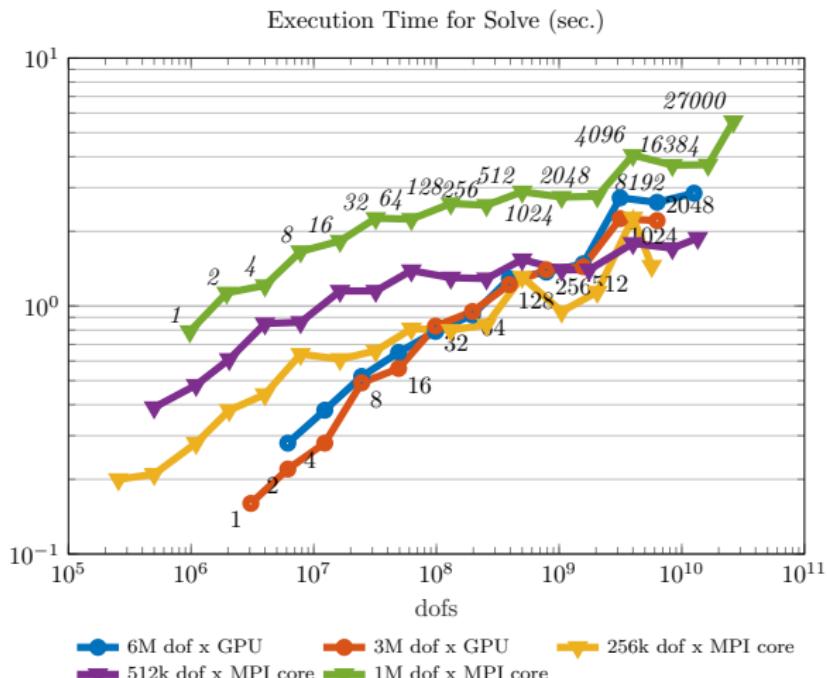
Let  $M$  be the spd (convergent)  $\ell_1$ -Jacobi smoother:

$$G = (I - M^{-1}A), \quad M = \text{diag}(M_{ii})_{i=1,\dots,n}$$
$$M_{ii} = a_{ii} + \sum_{j \neq i} |a_{ij}|$$

- Pros: simple and cheap to setup, only based on sparse matrix-vector product and local vector updates well suited for high-throughput SIMD processors
- Cons: larger approximation constant than block-Jacobi version of Gauss-Seidel (Hybrid GS); in our AMG setting the constant is  $\approx 4$  time larger for systems coming from 3D Poisson problem

# Parallel Smoothers

Some results on Piz Daint (ranked 43th): MPI-HGS vs MPI/CUDA-l1Jac



using GPUs saves up to  $\approx 50\%$

in solve time and energy consumption for 10 billion dofs

Polynomial accelerators (Adams et al. 2003, Kraus et al. 2012)

$$G = p_k((M')^{-1}A'), \text{ for } p_k(x) \in \Pi_k[x]$$

$$\text{s.t. } p_k(0) = 1 \text{ and } |p_k(x)| < 1 \text{ for } 0 < x \leq 1$$

Key issue: choose polynomials to optimize V-cycle approximation constant

# Parallel Smoothers

Polynomial accelerators (Adams et al. 2003, Kraus et al. 2012)

$$G = p_k((M')^{-1}A'), \text{ for } p_k(x) \in \Pi_k[x]$$

s.t.  $p_k(0) = 1$  and  $|p_k(x)| < 1$  for  $0 < x \leq 1$

Key issue: choose polynomials to optimize V-cycle approximation constant

Theorem (Lottes, 2023)

The V-cycle error propagation matrix has following bound:

$$\|E\|_A^2 \leq \max_I \frac{C^I}{C^I + (\gamma_k^I)^{-1}},$$

where  $C^I$  is the strong approximation property constant at the level  $I$  and

$$\gamma_k^I = \sup_{0 < \lambda \leq 1} \frac{\lambda p_k(\lambda)^2}{1 - p_k(\lambda)^2}$$

## Minimax problem

$$\gamma_k := \min_{p_k(x) \in \Pi_k} \max_{x \in (0,1]} \left| \frac{xp_k(x)^2}{1 - p_k(x)^2} \right|$$

s.t.  $p_k(0) = 1$  and  $|p_k(x)| < 1$  for  $0 < x \leq 1$

# Optimal V-cycle bound & Chebyshev polynomials

## Minimax problem

$$\gamma_k := \min_{p_k(x) \in \Pi_k} \max_{x \in (0,1]} \left| \frac{xp_k(x)^2}{1 - p_k(x)^2} \right|$$

s.t.  $p_k(0) = 1$  and  $|p_k(x)| < 1$  for  $0 < x \leq 1$

## Rewriting the minimax problem

$$\gamma_k = \min_{p_k(x) \in \Pi_k} \max_{x \in (0,1]} x \left| 1 - \frac{1}{1 - p_k(x)^2} \right|,$$

s.t.  $p_k(0) = 1$  and  $|p_k(x)| < 1$  for  $0 < x \leq 1$

# Optimal V-cycle bound & Chebyshev polynomials

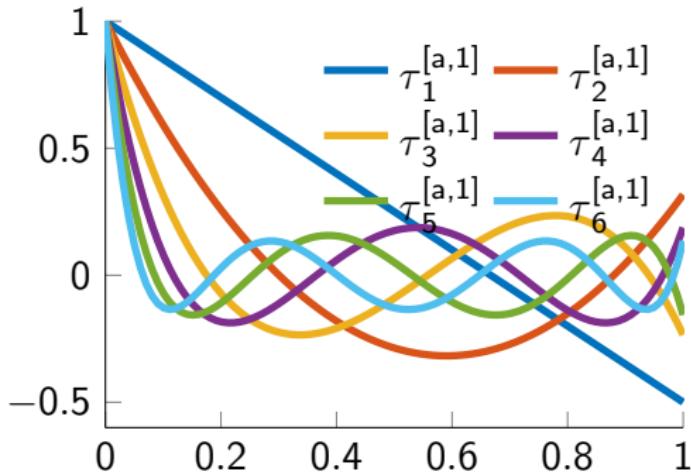
## Quasi-Optimal 1st-kind Chebyshev polynomials

$$\tau_k(x) = \frac{1}{2} \left[ (x + \sqrt{x^2 - 1})^k + (x - \sqrt{x^2 - 1})^k \right]$$

- $\tau_k(x)$  provides the optimal solution in the interval  $[a_k, 1]$ , for  $a_k \in (0, 1)$
- optimal values of  $a_k$  and corresponding  $\gamma_k$  can be numerically obtained by solving a scalar non-linear equation
- lower and upper bounds can be found for optimal  $a_k$  and  $\gamma_k$
- can be applied as a simple 3-terms recurrence requiring sparse matrix-vector products and vector updates

PD et al., *Optimal Polynomial Smoothers for Parallel AMG*, SIAM ALA, Paris, May 2024.

# Optimal V-cycle bound & Chebyshev polynomials

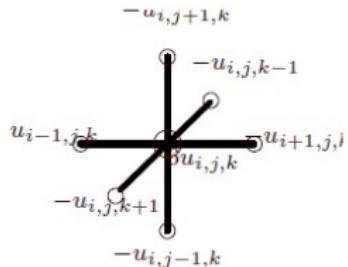


$$\mathbf{x}^{(k)} = \mathbf{x}^{(k-1)} + \mathbf{z}^{(k-1)}$$

$$\begin{cases} \mathbf{z}^{(0)} = \frac{2}{1+a^*} M^{-1}(\mathbf{b} - A\mathbf{x}^{(0)}) \quad \rho_0 = \frac{1-a^*}{1+a^*} \\ \rho_k = \left( \frac{2(1+a^*)}{1-a^*} - \rho_{k-1} \right)^{-1} \\ \mathbf{z}^{(k)} = \rho_k \rho_{k-1} \mathbf{z}^{(k-1)} + \frac{4\rho_k}{(1-a^*)} M^{-1}(\mathbf{b} - A\mathbf{x}^{(k)}) \end{cases}$$

# Test case: Poisson equation

$-\nabla^2 u = 1$  on unit cube, with DBC

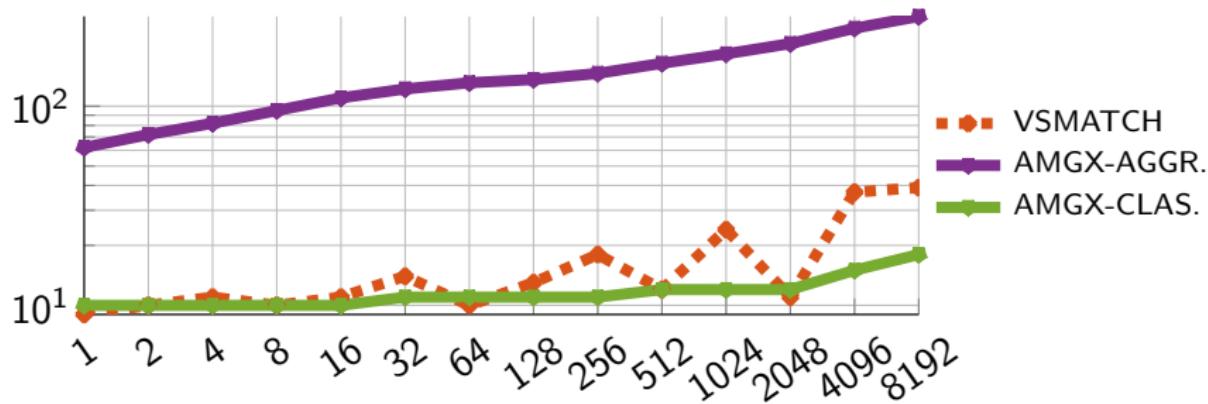


## Solver/preconditioner settings

- AMG as preconditioner of CG, stopped when  $\|\mathbf{r}^k\|_2/\|\mathbf{b}\|_2 \leq 10^{-6}$   
VSMATCH V-cycle for CMATCH-based AMG hierarchy with aggregates of max size 8, smoothed prolongators
- coarsest matrix size  $n_c \leq 200np$ , with  $np$  number of tasks (GPUs)
- 4  $\ell_1$ -Jacobi iterations as pre/post smoother; 40 iterations of PCG +  $\ell_1$ -Jacobi at the coarsest level.

Platform: Leonardo booster, BullSequana XH2000, Xeon Platinum 8358 32C 2.6GHz, NVIDIA A100 SXM4 64 GB, Quad-rail NVIDIA HDR100 Infiniband

# Iterations for Solve up to 65 billion dofs

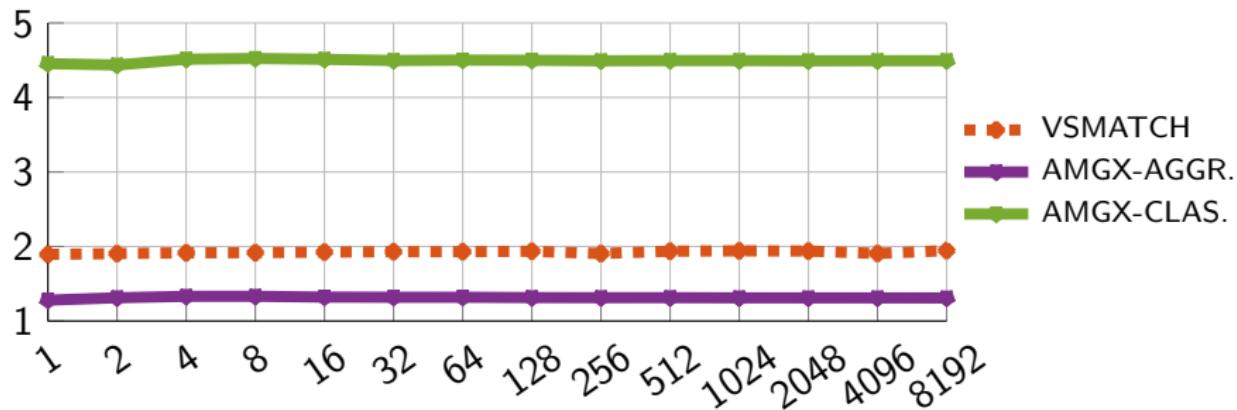


Comparison with Nvidia AmgX-PCG coupled with:

**AMGX-CLAS.** Ruge-Stüben coarsening, distance-2 interpolation, V-cycle, smoother  $\ell_1$ -Jacobi (4 sweeps),  $\ell_1$ -Jacobi (40 sweeps) coarsest solver

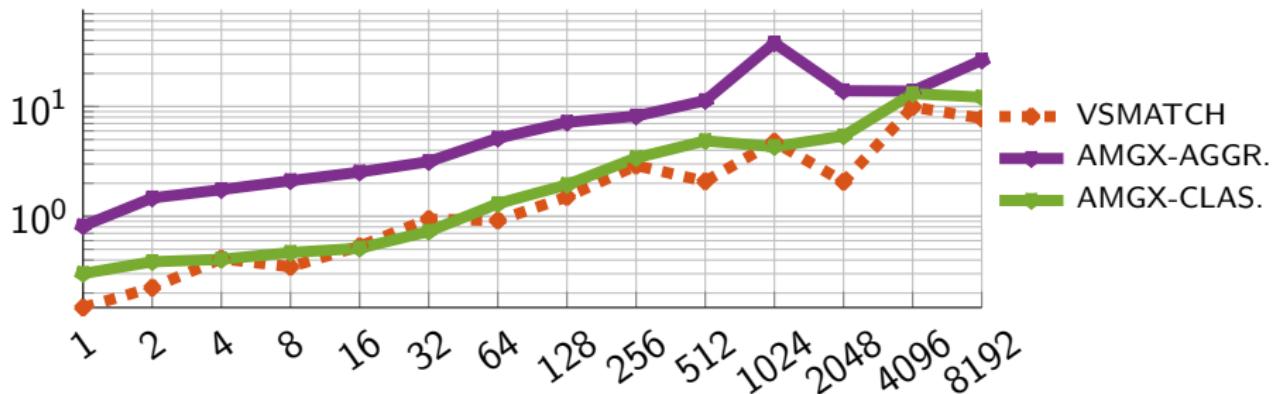
**AMGX AGGR.** aggregation by iterative parallel graph matching, aggregates of maximum size equal to 8, further algorithmic choices as in AMGX-CLAS

# Operator Complexity

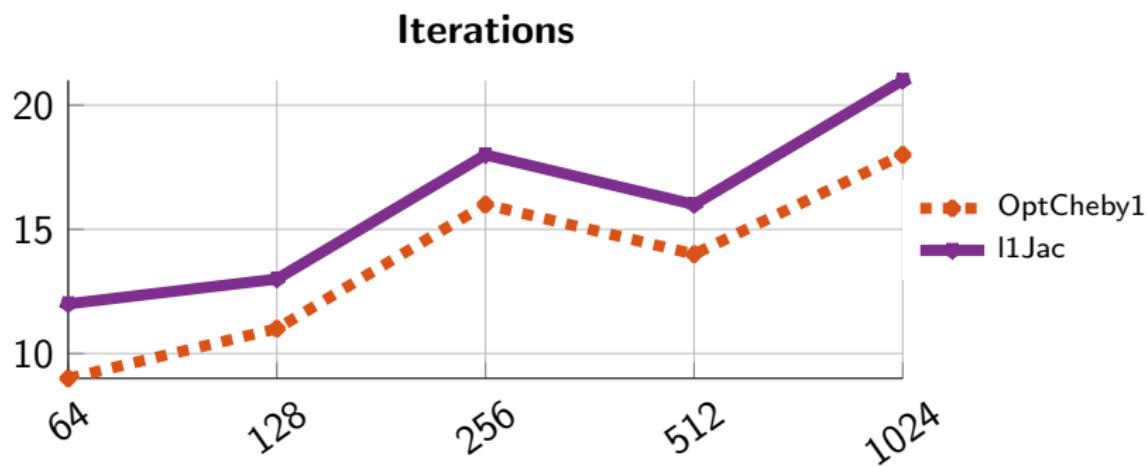


$$\frac{\sum_{l=0}^{\ell} \text{nnz}(A_l)}{\text{nnz}(A_0)}$$

# Solve Time (sec.)

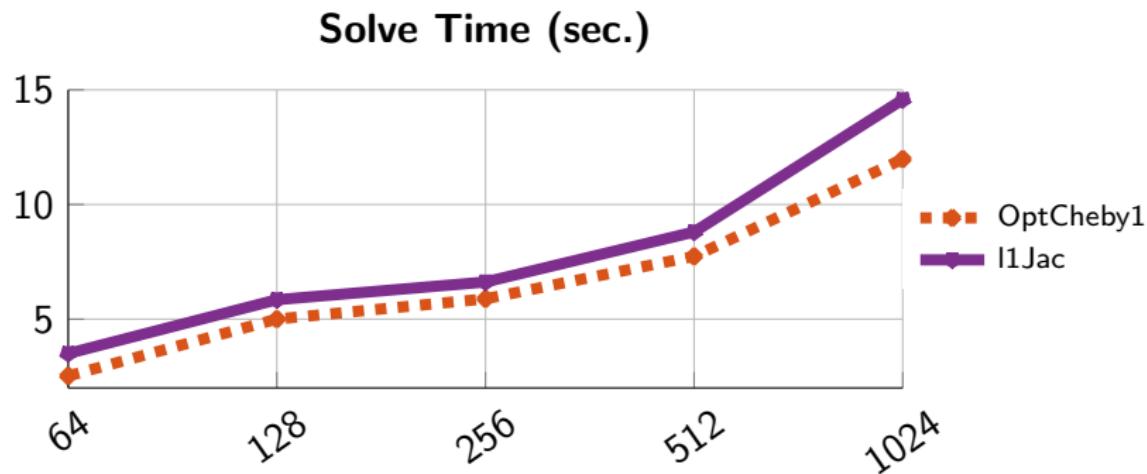


# VSMATCH: Opt-1st-kind Cheb.(deg 4) vs l1-Jac up (4 it)



17% iterations savings on 1024 GPUs for 6 billion dofs

# VSMATCH: Opt-1st-kind Cheb.(deg 4) vs l1-Jac up (4 it)



22% time savings on 1024 GPUs for 6 billion dofs

## Some concluding remarks

- PSCToolkit demonstrates benefits in solving benchmark systems up to 65 billion dofs on up to 8192 GPUs of the Leonardo supercomputer
- Applications to CFD for sustainable energy are in:
  - PD et al., *Alya towards Exascale: Algorithmic Scalability using PSCToolkit*, Journal of Supercomputing, Vol. 80, 2024.
  - PD et al., *Why diffusion-based preconditioning of Richards equation works: spectral analysis and computational experiments at very large scale*, Numerical Linear Algebra with Applications, Vol. 31, 1, 2024.
- Improvements in methods and software design are work in progress within different projects

Thanks for Your Attention

This work was performed with support of the European Union under HORIZON-EUROHPC-JU-2023-COE-01 agreement N. 101144014 — EoCoE-III and of the Italian Research Center on High-Performance Computing, Big Data and Quantum Computing (ICSC) funded by MUR - Next Generation EU (NGEU)

# Coarsening by Compatible Weighted Matching

Let  $\mathbf{w} \in \mathcal{R}^n$  smooth vector, let  $P_c \in \mathcal{R}^{n \times n_c}$  and  $P_f \in \mathcal{R}^{n \times n_f}$  be a **prolongator** and a **complementary prolongator**, such that:

$$\mathcal{R}^n = \text{Range}(P_c) \oplus^\perp \text{Range}(P_f), \quad n = n_c + n_f$$

$\mathbf{w} \in \text{Range}(P_c)$  : coarse space       $\text{Range}(P_f)$  : complementary space

$$[P_c, P_f]^T A [P_c, P_f] = \begin{pmatrix} P_c^T A P_c & P_c^T A P_f \\ P_f^T A P_c & P_f^T A P_f \end{pmatrix} = \begin{pmatrix} A_c & A_{cf} \\ A_{fc} & A_f \end{pmatrix}$$

$A_c$  : coarse matrix       $A_f$  : hierarchical complement

Compatible Relaxation (Falgout and Vassilevski, 2004)

Good estimate of the strong approximation constant for  $P_c$ :

$$\rho_f = \|I - M_f^{-1} A_f\|_{A_f} \text{ with } M_f = P_f^T M P_f$$

# Coarsening by Compatible Weighted Matching

Let  $\mathbf{w} \in \mathcal{R}^n$  smooth vector, let  $P_c \in \mathcal{R}^{n \times n_c}$  and  $P_f \in \mathcal{R}^{n \times n_f}$  be a **prolongator** and a **complementary prolongator**, such that:

$$\mathcal{R}^n = \text{Range}(P_c) \oplus^\perp \text{Range}(P_f), \quad n = n_c + n_f$$

$\mathbf{w} \in \text{Range}(P_c)$  : coarse space       $\text{Range}(P_f)$  : complementary space

$$[P_c, P_f]^T A [P_c, P_f] = \begin{pmatrix} P_c^T A P_c & P_c^T A P_f \\ P_f^T A P_c & P_f^T A P_f \end{pmatrix} = \begin{pmatrix} A_c & A_{cf} \\ A_{fc} & A_f \end{pmatrix}$$

$A_c$  : coarse matrix       $A_f$  : hierarchical complement

Compatible Relaxation (Falgout and Vassilevski, 2004)

Good estimate of the strong approximation constant for  $P_c$ :

$$\rho_f = \|I - M_f^{-1} A_f\|_{A_f} \text{ with } M_f = P_f^T M P_f$$

Our recipe

build  $P_c$  (and  $P_f$ ) by dofs **aggregation based on matching in the weighted (adjacency) graph of  $A$** , to make  $A_f$  as diagonally-dominant as possible

## Results: Time per Iteration

